

IRESH AGRAWAL

Data Engineer | Azure • Databricks • Snowflake

8707481550 • iresh.agrawal23@gmail.com • linkedin.com/in/iresh-agrawal23

PROFESSIONAL SUMMARY

Results-driven Data Engineer with 3.7+ years of experience architecting and optimizing cloud-native data platforms on Azure, Databricks, and Snowflake. Proven track record of delivering enterprise-scale data solutions that achieve 60%+ performance improvements and 100% SLA compliance with zero production defects. Expert in building scalable ETL/ELT pipelines using PySpark, implementing Data Vault 2.0 and Medallion Architecture, and driving successful Unity Catalog migrations for Fortune 500 retail and insurance clients. Microsoft Certified Azure Data Engineer (DP-203) with strong DevOps and CI/CD automation capabilities.

CORE TECHNICAL COMPETENCIES

Cloud Platforms:	Azure Databricks, Snowflake, Azure Data Factory (ADF), ADLS Gen2, Unity Catalog, AWS S3
Big Data & Processing:	Apache Spark, PySpark, Distributed Data Processing, Batch & Stream Processing, SQL Optimization
Data Architecture:	Data Vault 2.0 (Hubs, Links, Satellites), Medallion Architecture (Bronze-Silver-Gold), ETL/ELT Pipelines, Dimensional Modeling, Star Schema, Data Governance, Data Quality Frameworks
Programming:	SQL , Python, PySpark, Scala, Shell Scripting, XML
DevOps & CI/CD:	Jenkins, Azure DevOps Pipelines, GitHub, Git, Version Control, Automated Testing & Deployment
Tools & Analytics:	DBVisualizer, Power BI, JIRA, Rally, Putty, Microsoft Excel

PROFESSIONAL EXPERIENCE

EXL Services (via Clairvoyant) | Senior Executive – Software Engineer

Gurgaon/Pune, India | July 2025 – Present | Domain: Insurance

- Architected enterprise-grade Data Vault 2.0 models in Snowflake (Raw Vault, Business Vault, Clean Vault) across multiple insurance subject areas, establishing audit-compliant and analytics-ready data foundations supporting regulatory reporting and business intelligence
- Orchestrated 20+ database schema releases across Dev, QA, and Production environments using Jenkins CI/CD and GitHub version control, implementing XML-based master changelogs for controlled, auditable database change management
- Engineered standardized data ingestion frameworks using Shell scripting and SSH for automated Snowflake data loads from AWS S3, eliminating environment-specific dependencies and improving deployment reliability by 100%
- Led CI/CD implementation for two parallel development tracks (GOALD, NAVINS), establishing deployment standards that improved release predictability and reduced validation cycles by 35% using DBVisualizer
- Collaborated with 3+ client stakeholders, business analysts, and QA teams to translate functional requirements into production-ready Snowflake data models, ensuring alignment with business objectives across agile sprint cycles
- Mentored 3+ junior data engineers on Snowflake development best practices, Data Vault 2.0 modeling patterns, and enterprise release management processes, accelerating team capability development

Capgemini | Associate Consultant

Gurgaon & Bengaluru, India | November 2023 – June 2025 | Domain: FMCG/Retail

- Spearheaded Unity Catalog migration for 50TB+ enterprise data platform on Azure Databricks, implementing Medallion Architecture (Bronze-Silver-Gold layers) that achieved 62.5% reduction in pipeline execution time and 98% decrease in production defects
- Optimized complex data transformation pipelines processing 100+ Azure Data Factory sources using advanced PySpark and SQL techniques, increasing throughput by 40% and reducing data latency by 50% for time-sensitive retail analytics
- Managed end-to-end CI/CD pipeline lifecycle for 10+ production releases using Azure DevOps, achieving 100% on-time delivery with zero rollback incidents through rigorous testing and validation protocols
- Collaborated with onshore stakeholders and QA teams to validate 10+ production releases, ensuring data accuracy, integrity, and operational stability across global retail operations

Capgemini | Senior Analyst

Gurgaon, India | July 2022 – October 2023 | Domain: FMCG/Retail

- Resolved 95% of critical production incidents within SLA through systematic root cause analysis and implementation of permanent corrective actions, maintaining 99.9% platform uptime for mission-critical retail data systems
- Implemented comprehensive regression validation frameworks for ETL/ELT pipelines, reducing post-release failures by 40% and minimizing unplanned production downtime
- Developed Power BI dashboards and automated ETL workflows that reduced manual reporting effort by 40% and improved business intelligence delivery turnaround time by 45%
- Partnered with business users and QA teams to validate analytical outputs across multiple reporting dashboards, ensuring accuracy and reliability of business-critical metrics

Capgemini | Data Engineering Intern

Gurgaon, India | January 2022 – June 2022

- Developed Azure Data Factory pipelines for ingestion and transformation of structured and semi-structured datasets from diverse source systems
- Conducted comprehensive SQL-based data quality validation and testing across development and UAT environments

EDUCATION

Bachelor of Technology (B.Tech) in Electronics and Communication Engineering

Vellore Institute of Technology (VIT), Bhopal | July 2018 – September 2022

CERTIFICATIONS

- **Microsoft Certified: Azure Data Engineer Associate (DP-203)**
- Microsoft Certified: Azure Fundamentals (AZ-900)
- Databricks: Generative AI Fundamentals
- Capgemini: Jr. Data Engineer (Level 1 and 2) and Data Engineer Level 1

AWARDS & RECOGNITION

- **Exceeding Expectations Award (EXL Services)** – Recognized for architecting Snowflake Data Vault 2.0 models and streamlining CI/CD release management for enterprise insurance portfolios
- **Silver Star Award (Capgemini – 2x Recipient)** – Awarded twice for leading high-impact Unity Catalog migration and consistently achieving 100% SLA compliance in production support operations